

HW3: Explore and Exploit for Arm-Bandit Problem

Preparation

Create a Python base class for K-armed bandit algorithms.

Prompt:

Create a Python base class for K-armed bandit algorithms that:

- Serves as a foundation for various algorithms (Epsilon-Greedy, UCB, Softmax, Thompson Sampling)
- Has default parameters:
 - 10 arms
 - 1000 iterations
- Tracks the following metrics:
 - Count of selections for each arm
 - Cumulative average reward for each iteration
- Provides only the base class implementation
- Includes English comments throughout the code

Please provide the implementation with clear structure and documentation to facilitate easy extension with specific algorithms.

Code:

This is the base class for the K-armed bandit, which will be used in subsequent experiments.

```
In [11]: from abc import ABC, abstractmethod

import matplotlib.pyplot as plt
import numpy as np

class MultiArmedBandit(ABC):
    """
    Base class for Multi-Armed Bandit problems.
    This class provides the framework for implementing different bandit algorithms.
    """

    def __init__(self, n_arms=10, iterations=1000, true_rewards=None):
        """
        Initialize the Multi-Armed Bandit environment.

        Args:
            n_arms (int): Number of arms (actions)
            iterations (int): Number of iterations to run
        """
```

```

        true_rewards (list/array): Optional list of true reward means for ea
    """
    self.n_arms = n_arms
    self.iterations = iterations

    # Initialize true reward means for each arm if not provided
    if true_rewards is None:
        self.true_rewards = np.random.normal(0, 1, n_arms) # Random normal
    else:
        self.true_rewards = true_rewards

    # Tracking variables
    self.arm_counts = np.zeros(n_arms) # Count of times each arm was pulled
    self.rewards = np.zeros(n_arms) # Sum of rewards for each arm
    self.cumulative_rewards = np.zeros(iterations) # Cumulative average rew
    self.selected_arms = [] # History of selected arms

def reset(self):
    """Reset all tracking variables to start a new experiment"""
    self.arm_counts = np.zeros(self.n_arms)
    self.rewards = np.zeros(self.n_arms)
    self.cumulative_rewards = np.zeros(self.iterations)
    self.selected_arms = []

def pull_arm(self, arm_index):
    """
    Pull an arm and get its reward.

    Args:
        arm_index (int): The index of the arm to pull

    Returns:
        float: The reward from pulling the arm
    """
    # Generate reward with Gaussian noise around the true mean
    reward = np.random.normal(self.true_rewards[arm_index], 1)

    # Update counts and reward sums
    self.arm_counts[arm_index] += 1
    self.rewards[arm_index] += reward
    return reward

@abstractmethod
def select_arm(self):
    """
    Select which arm to pull next.
    This method should be implemented by subclasses with specific algorithms

    Returns:
        int: The index of the selected arm
    """
    pass

def run(self):
    """
    Run the bandit algorithm for the specified number of iterations.

    Returns:
        tuple: (cumulative_rewards, arm_counts, selected_arms)
    """

```

```

self.reset()
total_reward = 0

for t in range(self.iterations):
    # Select arm according to the algorithm
    arm = self.select_arm()
    self.selected_arms.append(arm)

    # Pull arm and observe reward
    reward = self.pull_arm(arm)

    # Update cumulative reward
    total_reward += reward
    self.cumulative_rewards[t] = total_reward / (t + 1)

return self.cumulative_rewards, self.arm_counts, self.selected_arms

def get_arm_mean_rewards(self):
    """
    Calculate the mean rewards for each arm based on observed rewards.

    Returns:
        array: Mean reward for each arm
    """
    mean_rewards = np.zeros(self.n_arms)
    for i in range(self.n_arms):
        if self.arm_counts[i] > 0:
            mean_rewards[i] = self.rewards[i] / self.arm_counts[i]
    return mean_rewards

def plot_results(self, color='blue'):
    """Plot the results of the experiment"""
    plt.figure(figsize=(15, 5))

    # Plot 1: Cumulative average reward over time
    plt.subplot(1, 2, 1)
    plt.plot(self.cumulative_rewards, color=color)
    plt.xlabel('Iteration')
    plt.ylabel('Cumulative Average Reward')
    plt.title('Performance Over Time')

    # Plot 2: Arm selection counts
    plt.subplot(1, 2, 2)
    plt.bar(range(self.n_arms), self.arm_counts, color=color)
    plt.xlabel('Arm')
    plt.ylabel('Number of Pulls')
    plt.title('Arm Selection Frequency')

    plt.tight_layout()
    plt.show()

```

Experiment

1. Epsilon-Greedy

Prompt:

Generate Latex code for the Epsilon-Greedy algorithm.

Implement the Epsilon-Greedy algorithm based on the provided base class. And then run the algorithm with the following parameters:

- 10 arms
- 1000 iterations
- Epsilon = 0.3

Plot the average reward over time and the number of selections for each arm.

Latex:

Algorithm 4 Thompson Sampling (Beta-Binomial for Binary Rewards)

- 1: **Input:** Number of arms K , time horizon T
 - 2: Initialize: $\alpha(a) \leftarrow 1, \beta(a) \leftarrow 1$ for all arms $a \in \{1, \dots, K\}$ \triangleright Prior for Beta distribution
 - 3: **for** $t = 1$ to T **do**
 - 4: **for** each arm $a \in \{1, \dots, K\}$ **do**
 - 5: Sample $\theta(a) \sim \text{Beta}(\alpha(a), \beta(a))$
 - 6: **end for**
 - 7: Choose arm $a_t = \arg \max_a \theta(a)$
 - 8: Observe reward $r_t \in \{0, 1\}$ from arm a_t \triangleright Assumes binary rewards
 - 9: Update $\alpha(a_t) \leftarrow \alpha(a_t) + r_t$
 - 10: Update $\beta(a_t) \leftarrow \beta(a_t) + (1 - r_t)$
 - 11: **end for**
-

Code:

```
In [12]: class EpsilonGreedy(MultiArmedBandit):
    """
    Implementation of Epsilon-Greedy algorithm for the Multi-Armed Bandit problem

    The algorithm selects the best arm with probability (1-epsilon) and
    explores a random arm with probability epsilon.
    """

    def __init__(self, n_arms=10, iterations=1000, epsilon=0.1, true_rewards=None):
        """
        Initialize the Epsilon-Greedy algorithm.

        Args:
            n_arms (int): Number of arms (actions)
            iterations (int): Number of iterations to run
            epsilon (float): Exploration parameter between 0 and 1
            true_rewards (list/array): Optional list of true reward means for each arm
        """
        super().__init__(n_arms, iterations, true_rewards)
        self.epsilon = epsilon

    def select_arm(self):
        """
        Select an arm using the Epsilon-Greedy strategy.

        With probability (1-epsilon), select the arm with the highest estimated
```

With probability `epsilon`, select a random arm.

Returns:

int: The index of the selected arm
"""

Exploration: select a random arm with probability epsilon

```
if np.random.random() < self.epsilon:  
    return np.random.randint(self.n_arms)
```

Exploitation: select the best arm with probability (1-epsilon)

else:

For arms that haven't been tried yet, assign them a high value to
estimated_rewards = np.zeros(self.n_arms)

```
for i in range(self.n_arms):
```

```
    if self.arm_counts[i] > 0:
```

```
        estimated_rewards[i] = self.rewards[i] / self.arm_counts[i]
```


```
    else:
```

```
        estimated_rewards[i] = float('inf') # Ensure untried arms a
```

Return the arm with the highest estimated reward

```
return np.argmax(estimated_rewards)
```

```
In [13]: eg = EpsilonGreedy(n_arms=10, iterations=1000, epsilon=0.3)  
eg_cumulative_rewards, eg_arm_counts, eg_selected_arms = eg.run()  
eg.plot_results(color='blue')
```



2. UCB (Upper Confidence Bound)

Prompt:

Generate Latex code for the UCB algorithm.

Implement the UCB algorithm based on the provided base class. And then run the algorithm with the following parameters:

- 10 arms
- 1000 iterations
- $C = 2$

Plot the average reward over time and the number of selections for each arm.

Latex:

Algorithm 2 UCB

Input: Number of arms K , iterations T , exploration parameter $c > 0$

Initialize: $Q(a) \leftarrow 0$, $N(a) \leftarrow 0$ for all arms $a \in \{1, \dots, K\}$, $t \leftarrow 0$

for $t = 1$ to T **do**

if there exists arm a with $N(a) = 0$ **then**

 Select arm $a_t \leftarrow a$

else

 Select arm $a_t \leftarrow \arg \max_a \left[Q(a) + c \sqrt{\frac{\ln t}{N(a)}} \right]$

end if

 Observe reward r_t from arm a_t

 Update: $N(a_t) \leftarrow N(a_t) + 1$

 Update: $Q(a_t) \leftarrow Q(a_t) + \frac{1}{N(a_t)}(r_t - Q(a_t))$

end for

Code:

```
In [14]: import math

class UCB(MultiArmedBandit):
    """
    Implementation of Upper Confidence Bound (UCB) algorithm for the Multi-Armed
    Bandit problem.

    UCB selects actions according to:
    UCB_i = Q_i + c * sqrt(ln(t)/N_i)

    where:
    - Q_i: estimated reward mean for arm i
    - c: exploration parameter
    - t: total number of rounds so far
    - N_i: number of times arm i has been pulled
    """

    def __init__(self, n_arms=10, iterations=1000, c=2.0, true_rewards=None):
        """
        Initialize the UCB algorithm.

        Args:
            n_arms (int): Number of arms (actions)
            iterations (int): Number of iterations to run
            c (float): Exploration parameter controlling the confidence bounds
            true_rewards (list/array): Optional list of true reward means for ea
        """
        super().__init__(n_arms, iterations, true_rewards)
        self.c = c
        self.t = 0 # Total number of rounds

    def reset(self):
        """Reset all tracking variables to start a new experiment"""
        super().reset()
        self.t = 0

    def select_arm(self):
        """
        Select an arm using the UCB strategy.
        """
```

Calculate the UCB value for each arm and select the arm with the highest
For arms that haven't been tried yet, assign them a high value to ensure

Returns:

```
    int: The index of the selected arm
"""
self.t += 1

# First, make sure each arm is tried at least once
for i in range(self.n_arms):
    if self.arm_counts[i] == 0:
        return i


# Calculate UCB value for each arm
ucb_values = np.zeros(self.n_arms)
for i in range(self.n_arms):
    # Estimated mean reward
    mean_reward = self.rewards[i] / self.arm_counts[i]

    # Exploration bonus
    exploration_bonus = self.c * math.sqrt(math.log(self.t) / self.arm_c

    # UCB value
    ucb_values[i] = mean_reward + exploration_bonus

# Return the arm with the highest UCB value
return np.argmax(ucb_values)
```

```
In [15]: ucb = UCB(n_arms=10, iterations=1000)
ucb_cumulative_rewards, ucb_arm_counts, ucb_selected_arms = ucb.run()
ucb.plot_results(color='green')
```



3. Softmax

Prompt:

Generate Latex code for the Softmax algorithm.

Implement the Softmax algorithm based on the provided base class. And then run the algorithm with the following parameters:

- 10 arms
- 1000 iterations
- Temperature = 0.2

Latex:

Algorithm 3 Softmax

Input: Number of arms K , iterations T , temperature $\tau > 0$
Initialize: $Q(a) \leftarrow 0$, $N(a) \leftarrow 0$ for all arms $a \in \{1, \dots, K\}$
for $t = 1$ to T **do**
 if there exists arm a with $N(a) = 0$ **then**
 Select arm $a_t \leftarrow a$
 else
 Compute probabilities: $p(a) \leftarrow \frac{\exp(Q(a)/\tau)}{\sum_{b=1}^K \exp(Q(b)/\tau)}$ for all a
 Select arm $a_t \sim \text{Categorical}(p(1), \dots, p(K))$
 end if
 Observe reward r_t from arm a_t
 Update: $N(a_t) \leftarrow N(a_t) + 1$
 Update: $Q(a_t) \leftarrow Q(a_t) + \frac{1}{N(a_t)}(r_t - Q(a_t))$
end for

```
In [16]: class Softmax(MultiArmedBandit):
    """
    Implementation of Softmax algorithm for the Multi-Armed Bandit problem.

    The Softmax algorithm selects arms with probability proportional to their
    exponentially weighted average rewards, controlled by a temperature parameter.
    Higher temperature leads to more exploration (more uniform probabilities).
    Lower temperature leads to more exploitation (higher probability for better
    """

    def __init__(self, n_arms=10, iterations=1000, temperature=0.2, true_rewards=None):
        """
        Initialize the Softmax algorithm.

        Args:
            n_arms (int): Number of arms (actions)
            iterations (int): Number of iterations to run
            temperature (float): Temperature parameter controlling exploration
                                Higher value = more exploration
                                Lower value = more exploitation
            true_rewards (list/array): Optional list of true reward means for each arm
        """
        super().__init__(n_arms, iterations, true_rewards)
        self.temperature = temperature

    def select_arm(self):
        """
        Select an arm using the Softmax strategy.

        Converts the estimated mean rewards of each arm into a probability distribution
        using the softmax function, then samples an arm according to that distribution.

        Returns:
            int: The index of the selected arm
        """
        # First, ensure each arm is tried at least once
        for i in range(self.n_arms):
            if self.arm_counts[i] == 0:
                return i
```



```

# Calculate estimated mean rewards
estimated_rewards = np.zeros(self.n_arms)
for i in range(self.n_arms):
    estimated_rewards[i] = self.rewards[i] / self.arm_counts[i]

# Apply softmax function to convert rewards to probabilities
# To prevent overflow, subtract the maximum reward before exponentiating
max_reward = np.max(estimated_rewards)
exp_rewards = np.exp((estimated_rewards - max_reward) / self.temperature)
probabilities = exp_rewards / np.sum(exp_rewards)


# Sample an arm according to the calculated probabilities
return np.random.choice(self.n_arms, p=probabilities)

```

```

In [17]: softmax = Softmax(n_arms=10, iterations=1000)
softmax_cumulative_rewards, softmax_arm_counts, softmax_selected_arms = softmax.
softmax.plot_results(color='red')

```



4. Thompson Sampling

Prompt:

Generate Latex code for the Thompson Sampling algorithm.

Implement the Thompson Sampling algorithm based on the provided base class. And then run the algorithm with the following parameters:

- 10 arms
- 1000 iterations
- Beta distribution parameters: $\alpha = 1$, $\beta = 1$

Latex:

Algorithm 4 Thompson Sampling (Beta Distribution)

Input: Number of arms K , iterations T , prior parameters $\alpha_0, \beta_0 > 0$

Initialize: $\alpha(a) \leftarrow \alpha_0, \beta(a) \leftarrow \beta_0$ for all arms $a \in \{1, \dots, K\}$

for $t = 1$ to T **do**

for each arm $a \in \{1, \dots, K\}$ **do**

 Sample $\theta(a) \sim \text{Beta}(\alpha(a), \beta(a))$

end for

 Select arm $a_t \leftarrow \arg \max_a \theta(a)$

 Observe reward r_t from arm a_t

 Scale reward: $r'_t \leftarrow \frac{r_t - \min(r)}{\max(r) - \min(r)}$

▷ If $\max(r) \neq \min(r)$

 Update: $\alpha(a_t) \leftarrow \alpha(a_t) + r'_t$

 Update: $\beta(a_t) \leftarrow \beta(a_t) + (1 - r'_t)$

end for

Code:

```
In [18]: class ThompsonSampling(MultiArmedBandit):
        """
        Implementation of Thompson Sampling algorithm for the Multi-Armed Bandit problem.

        Thompson Sampling uses Bayesian approach by maintaining a probability distribution
        over the reward for each arm. For each action, it samples from these distributions
        and selects the arm with the highest sampled value.

        This implementation uses Beta distributions to model the rewards for each arm, which
        is suitable for rewards in the [0,1] range. For non-binary rewards, the implementation
        scales and shifts rewards to the [0,1] interval.
        """

        def __init__(self, n_arms=10, iterations=1000, alpha=1.0, beta=1.0, true_rewards=None):
            """
            Initialize the Thompson Sampling algorithm.

            Args:
                n_arms (int): Number of arms (actions)
                iterations (int): Number of iterations to run
                alpha (float): Initial alpha parameter for Beta distribution
                beta (float): Initial beta parameter for Beta distribution
                true_rewards (list/array): Optional list of true reward means for each arm
            """
            super().__init__(n_arms, iterations, true_rewards)
            self.alpha = np.ones(n_arms) * alpha # Success counts for each arm (prior)
            self.beta = np.ones(n_arms) * beta # Failure counts for each arm (prior)

            # For scaling rewards
            self.min_reward = float('-inf')
            self.max_reward = float('inf')

        def reset(self):
            """Reset all tracking variables to start a new experiment"""
            super().reset()
            self.alpha = np.ones(self.n_arms) * 1.0
            self.beta = np.ones(self.n_arms) * 1.0
            self.min_reward = float('-inf')
            self.max_reward = float('inf')
```

```

def pull_arm(self, arm_index):
    """
    Pull an arm and get its reward, updating the Beta distribution parameter

    Args:
        arm_index (int): The index of the arm to pull

    Returns:
        float: The reward from pulling the arm
    """
    # Get reward from the base class method
    reward = super().pull_arm(arm_index)

    # Update min and max for scaling
    self.min_reward = min(self.min_reward, reward)
    self.max_reward = max(self.max_reward, reward)

    # Scale reward to [0, 1] if we have seen multiple rewards
    if self.min_reward < self.max_reward:
        scaled_reward = (reward - self.min_reward) / (self.max_reward - self.min_reward)
    else:
        scaled_reward = 0.5 # Default if all rewards are the same

    # Update Beta distribution parameters
    self.alpha[arm_index] += scaled_reward
    self.beta[arm_index] += (1.0 - scaled_reward)

    return reward

def select_arm(self):
    """
    Select an arm using the Thompson Sampling strategy.

    Sample a value from the Beta distribution for each arm and select the arm
    with the highest sampled value.

    Returns:
        int: The index of the selected arm
    """
    # Sample a value from Beta distribution for each arm
    samples = np.zeros(self.n_arms)
    for i in range(self.n_arms):
        samples[i] = np.random.beta(self.alpha[i], self.beta[i])


    # Return the arm with the highest sampled value
    return np.argmax(samples)

```

```

In [19]: ts = ThompsonSampling(n_arms=10, iterations=1000, alpha=1.0, beta=1.0)
         ts_cumulative_rewards, ts_arm_counts, ts_selected_arms = ts.run()
         ts.plot_results("orange")

```



Experiment Results

Prompt:

Compare the cumulative rewards for each algorithm and plot the results.

Conduct space and time complexity analysis, explaining the performance and differences of each algorithm.


Code:

```
In [20]: # Compare the cumulative_rewards of each algorithm
plt.figure(figsize=(10, 6))

# Plot the cumulative_rewards for each algorithm
plt.plot(eg_cumulative_rewards, label='Epsilon-Greedy', color='blue')
plt.plot(ucb_cumulative_rewards, label='UCB', color='green')
plt.plot(softmax_cumulative_rewards, label='Softmax', color='red')
plt.plot(ts_cumulative_rewards, label='Thompson Sampling', color='orange')

# Add Legend, Labels, and title
plt.xlabel('Iterations')
plt.ylabel('Cumulative Average Reward')
plt.title('Comparison of Cumulative Rewards for Each Algorithm')
plt.legend()
plt.grid(True)

# Display the plot
plt.show()
```



Explanation:

Spatial (Performance) Analysis

- **Softmax** achieves the highest cumulative reward (~2.0), demonstrating superior performance
- **Thompson Sampling** performs second-best with a final reward of ~1.8
- **UCB (Upper Confidence Bound)** shows moderate performance (~1.0)
- **Epsilon-Greedy** has the lowest performance (~0.5)

Temporal (Convergence) Analysis

- **Softmax** shows early spikes but converges quickly (by ~200 iterations)
- **Thompson Sampling** has initial volatility but stabilizes around iteration 300
- **UCB** shows consistent, gradual improvement throughout the experiment
- **Epsilon-Greedy** starts with negative rewards, slowly improves, and plateaus earliest

Algorithm Comparison

Algorithm	Performance	Convergence	Key Characteristics
Softmax	Excellent (~2.0)	Fast	Uses probability distribution based on estimated values; temperature parameter controls exploration
Thompson Sampling	Good (~1.8)	Medium	Bayesian approach; samples from posterior distributions; naturally balances exploration/exploitation
UCB	Fair (~1.0)	Slow but steady	Selects arms based on upper confidence bounds; exploration decreases with more pulls

Algorithm	Performance	Convergence	Key Characteristics
Epsilon-Greedy	Poor (~0.5)	Slow	Simple implementation; random exploration with probability ϵ ; otherwise exploitation

Key Differences

The algorithms fundamentally differ in their exploration-exploitation strategies:

- **Softmax** dynamically adjusts action probabilities based on estimated rewards, allowing for nuanced exploration
- **Thompson Sampling** uses Bayesian principles to naturally balance exploration and exploitation through posterior sampling
- **UCB** systematically explores by considering uncertainty in reward estimates, favoring actions with higher potential
- **Epsilon-Greedy** uses a fixed exploration rate, leading to potentially suboptimal exploration patterns

The results suggest that in this environment, approaches with adaptive exploration strategies (Softmax, Thompson Sampling) significantly outperform simpler methods like Epsilon-Greedy, with Softmax's probability-based selection showing the best overall performance.